EBR/QBR Service Blueprint — all stages, in the map

Sourced from `briefs/ebr-process-map-2026-07-28.md` (Maps 1–6) + `briefs/context/qbr-ebr.context.md` §2/§6 (the real 15-tab template). Prep-Assemble is no longer one collapsed column — "prep the QBR" is 8 separately-sourced data pulls, each named, each tied to a specific tool, the same way a real experience map never lets "research the trip" stand in for "look up time tables" / "research hotels" / "compare airfare" as if they were one action.

The model: a service blueprint is an experience map with the back of the house added. Rows 1–5 are the customer's experience map (PICNIC bands). Below the Line of Visibility, rows 6–8 are the operational rows a pure experience map never shows. 26 columns: Entice · Enter · Schedule · 8 named Prep-Assemble pulls · Prep-Judge · the 12 live meeting moments (Stefanie's own tab order) · Follow-up · Extend.

Craig
Owner/Partner — buyer/champion
Went to Empower, brought the framework home. Not a swimlane actor in a single cycle — he's why the practice exists.

Vanessa Chart
Client Delivery Lead — the user
4–5h/EBR, hit her own ceiling at 4/week. Deliberately not a technician.

Head of Technical Overview
Hidden cost center
Pulled off the service desk to build the risk review.

Client Executive
The customer in this blueprint
Not interested in the back end; arrives guarded if the MSP's gone quiet.

What this map is built from — read this before quoting anything off it.
Small-n qualitative research only. Two sources: one taught process (Stefanie Hammond, N-able Head Nerd — a recorded training session plus her real 15-tab EBR template) and **one** MSP operator's account of running that process for a year (Vanessa Chart, Technica Solutions — n=1, a vendor-selected reference customer). Both are vendor content. That makes this map **directional, not projectable**: it tells you what the work actually looks like and where it hurts. It does not tell you how common any of it is, and no figure on this page should be read as a market number.
The receipts: `briefs/ebr-process-map-2026-07-28.md` (Maps 1–6, timestamped quotes) · `briefs/context/qbr-ebr.context.md` (the standing spec, fully footnoted) · `scout_input/stephanie-hammond-ebr-sources-2026-07-28.md` (raw source list) · `Executive Business Review Template 1.xlsx` (Stefanie's actual template). Jargon index is at the bottom of this page.

Tab classification (Maister's leverage-ratio distinction, `qbr-ebr.context.md` §6): `production-line` pullable/automatable · `hybrid` entries pullable, synthesis is judgment · `brains-work – protect this` live conversation only · `unclassified` not sorted in the original pass, marked honestly.

Bottom row — **our own build, checked against the map** (`acp-core-main`, read 2026-07-29): `BUILT` code exists · `SPECCED` in a spec, no code · `ADJACENT` something related exists but doesn't do this · `ABSENT` nothing.

	ENTICE	ENTER	SCHEDULE	PREP			
SERVICE MOMENT	pre-sale, step 1	onboarding, steps 2–4	booked live, prior cycle's close	Assemble — 8 separate data pulls, not one step			
WHICH PULL				**1. Warranty/EOL** Is this hardware still covered by the maker's warranty, and when does the maker stop patching it at all ("End-of-Life")? Unsupported hardware can't be fixed even once you know it's broken.	**2. License utilization** Of the software seats the client pays for every month, which ones has literally nobody logged into in 3–4 months?	**3. User/staff reconcile** Does the list of people who can log in match the list of people who still work there? Catches ex-employees who still have access and new hires who don't yet.	**4. Ticket volume + SLA** How many support requests came in, and did the MSP answer/fix them inside the response time it contractually promised (its "SLA")?
ACTIONS (CUSTOMER)	Prospect meets AM before deal closes; forms early expectations.	Attends kickoff, receives welcome gift, joins check-ins.	Agrees to and confirms the next EBR date, live, before leaving the room.	*None yet, across all 8 pulls — backstage-only, client not present.*			
PROCESSES, TOOLS, TECH	Sales conversation, in-person/call.	Kickoff meeting, welcome gift, check-in calls.	Verbal, in-room — no scheduling tool observed.	*No customer touchpoint at any of the 8 pulls — each pull's specific tool is in Support Processes below.*			
THINKING & FEELING	*[?]*	*[?]*	*Implied momentum if booking is easy.*	*n/a across all 8 pulls — client not present.*			
STRESSORS	*[?]*	**No structure/no framework** — "they don't know how to structure them."	No system backs the live commit up — if it doesn't happen, cadence discipline resets to whatever's convenient later.	**Resource contention across all 8 pulls** — one person (Head of Technical Overview) doing all of them, pulled off paying service-desk work for the whole batch, not per-pull.			
OPPORTUNITIES (BUILD IMPLICATION)	*[?] open — no material.*	The kit's structural-templating value answers Rib 1 directly — if it's the right structure (Musts 1–8).	**Build note:** the live verbal commit may be load-bearing for show-rate; a calendar-bot replacing it is untested risk.	⚙ **Direct API pull**, no judgment. Feeds Hardware/Software Asset Review + Current Climate.	⚙ **Kit the detection** — inactive-login query against login logs is mechanical; *what to do about it* is judgment, feeds the Business Priorities conversation, not this pull.	⚙ **Kit the diff** — two-list comparison (AD/IdP vs. HR roster) flags joiners/leavers; provisioning follow-up is a separate action, not this pull.	⚙ **Kit this** — a standard PSA report export. Feeds Current Climate.
— LINE OF VISIBILITY — only the business-value translation and the client-facing extract cross here —							
STAFF ACTIONS — FRONTSTAGE —	**AM**: "pre-frames" the relationship.	**AM**: runs kickoff, support-process education.	**Vanessa**: books next EBR live, every time.	*Invisible across all 8 pulls — assembly is backstage-only.*			
STAFF ACTIONS — BACKSTAGE —	*[?] who preps the AM's pitch — unknown.*	*[?] who preps kickoff materials — unknown.*	*[?] which cadence tier (A/B/C) applies — assumed known, not confirmed as a system lookup.*	⚙ **Head of Tech. Overview**: pulls hardware warranty status + EOL dates.	⚙ **Head of Tech. Overview**: checks license utilization — who hasn't logged in 3–4mo.	⚙ **Head of Tech. Overview**: reconciles user/employee list.	⚙ **Head of Tech. Overview**: pulls ticket volume + SLA attainment for the period.
SUPPORT PROCESSES	*[?] none named.*	*[?] none named.*	*[?] no scheduling tool named.*	`Asset/RMM inventory system (warranty lookup, EOL database).`	`License management tool / M365 admin center / IdP login logs.`	`Active Directory or IdP vs. HR roster (manual or API cross-check).`	`PSA (Halo) reporting.`
— BELOW THIS LINE IS NOT THE MSP'S SERVICE. IT IS OUR OWN BUILD, CHECKED AGAINST IT (acp-core-main, read 2026-07-29) —							
ACP TODAY (IN THE BUILD)	ABSENT Nothing. No pre-sale surface. –	ADJACENT There *is* an onboarding spec — but it onboards **the MSP into ACP**, not the MSP's client into the MSP. Same word, inverted direction. `Intent Specs/FSN-82692`	ADJACENT **ScheduleService** can fire a coworker on a cron/interval and is unit-tested — but it is called from **no API route**. Built infrastructure, no product surface. `app/runs/temporal/schedules.py`	ABSENT No warranty or EOL concept anywhere in code or specs.	ABSENT No license/seat/login-inactivity check. (Value Ledger names `license_savings` as a future secondary-value idea — a savings estimate, not a dormant-seat query. `specs/003-value-ledger/spec.md:186`)	ABSENT No AD/IdP/HR-roster reconcile.	ADJACENT PSA exists only as narrative: a "Ticket Triage & Dispatch" coworker story. ConnectWise/Autotask/Halo are *named* and marked **"pending spec — not started."** `Intent Specs/FSN-82692:437`

	PREP				
SERVICE MOMENT					Judge ★
WHICH PULL	**5. Patch + backup** Are systems up to date on security fixes, and are backups running — and, separately, has anyone tested that a backup can actually be *restored*? "Job succeeded" and "restorable" are different claims.	**6. Phish-sim + training** Did staff pass the fake phishing emails sent to test them ("phish-sim"), and did they finish required security-awareness training?	**7. Tenant config vs. baseline** Is the client's cloud environment (their "tenant" — e.g. their own Microsoft 365) still set up the way the MSP recommends (the MSP's own "baseline"), or has it drifted?	**8. Compliance checklist** Is the client meeting whatever regulation applies to them — HIPAA for healthcare, PCI for card payments, etc.? Which standards apply per client isn't named in source.	
ACTIONS (CUSTOMER)					*None yet.*
PROCESSES, TOOLS, TECH					—
THINKING & FEELING					*n/a — client not present*
STRESSORS					Operator discomfort in an exec room — no product answer exists.
OPPORTUNITIES (BUILD IMPLICATION)	⚙ **Kit this** — direct pull; **explicitly flag overdue restore tests**, not just backup-job-success — "ran" and "verified restorable" are different claims and the source only confirms the former reliably.	⚙ **Kit this** — a completion-rate export from the awareness platform.	**Hybrid** — the config diff is pullable, but "vs. OUR standard" means the baseline itself is MSP-authored judgment, not universal. Don't kit the baseline definition, only the diff against it.	**Hybrid** — checklist items pullable if the standard is codified; the narrative "are we compliant enough" stays judgment.	Draft the business-consequence translation. **Do not** build a "documentation shields you legally" claim.
STAFF ACTIONS — FRONTSTAGE —					*Invisible — judgment is backstage-only.*
STAFF ACTIONS — BACKSTAGE —	⚙ **Head of Tech. Overview**: checks patch compliance + backup job success/restore-test status.	⚙ **Head of Tech. Overview**: reviews phish-sim results + awareness-training completion.	★ **Head of Tech. Overview**: checks tenant security config against the MSP's own baseline (baseline itself is MSP-authored).	★ **Head of Tech. Overview**: pulls status against a named compliance-standard checklist.	★ **Vanessa**: judges gap vs. standard, translates to business consequence, sets roadmap priority — the real constraint.
SUPPORT PROCESSES	`RMM (patch status) + backup-monitoring tool (job success + restore-test log).`	`Security-awareness / phish-sim platform.`	`Tenant security-posture tool, compared against an internally-defined baseline.`	`Compliance/GRC tool, or a manual checklist if none exists.`	*None named — the one moment with no tool at all.*
ACP TODAY (IN THE BUILD)	ADJACENT A seeded demo "Patch Management Expert" persona and a `backup.restore` string in one unit test. No RMM, no backup-monitoring connector. `scripts/seed_experts.py:37`	ABSENT No phish-sim or awareness-training surface.	ABSENT No posture tool, no baseline-comparison mechanism.	ABSENT No GRC or compliance-checklist surface.	ABSENT **Nothing — and this is the one step the MSP has no tool for either.** Absent on both sides of the map.

	ENGAGE — THE MEETING ITSELF (STEFANIE'S TAB ORDER)				
SERVICE MOMENT	Meeting Kick Off unclass.	Current Climate prod-line	Risk & Security Review hybrid	Compliancy Review hybrid	Business Priorities & Goals brains
WHICH PULL					
ACTIONS (CUSTOMER)	Arrives, hears purpose/agenda recap.	Reviews snapshot of current environment.	Reviews risk register; asks questions.	Reviews compliance standing.	Discusses actual business goals/priorities; gives or withholds consensus.
PROCESSES, TOOLS, TECH	**The meeting itself**, in-person, one continuous session — one instance ran through lunch. (single touchpoint spans all 12 live moments, per the guide's "one touchpoint per service moment" convention read at the meeting level, not the tab level.)				
THINKING & FEELING	*[?]*	*[?]*	"Not really interested in the back end." If translation succeeds, feels informed, not overwhelmed.	*[?]*	Engagement here is what makes the later referral ask land.
STRESSORS	*[?]*	*[?]*	**The load-bearing moment.** "Can't prepare 20 min before you walk in." If discomfort wins here: raw data shown instead of translation, meeting reads tactical.	*[?]*	If this collapses into tech/stats recitation, client sees no value, declines next.
OPPORTUNITIES (BUILD IMPLICATION)	*[?] no material.*	Low-risk to auto-generate this tab's content once the 8 pulls feed it.	Confirm-gate any Unverified claim before it's spoken. Restate any prep-time claim as a ratio, never an absolute.	Entries pullable; narrative synthesis stays judgment.	**Do not template.** Brains-work — protect it. A product surface should support Vanessa's live judgment, not replace it.
STAFF ACTIONS — FRONTSTAGE —	**Vanessa**: opens meeting, states purpose.	**Vanessa**: presents snapshot.	**Vanessa** (NOT a technician): delivers pre-drafted translation, parks tech Qs.	**Vanessa**: presents compliance standing.	**Vanessa**: reads the room, builds consensus, sets priority live.
STAFF ACTIONS — BACKSTAGE —	*Already completed in Prep — invisible during the live meeting.*				
SUPPORT PROCESSES	*n/a during the live meeting — all support processes are Prep-side.*				
ACP TODAY (IN THE BUILD)	ABSENT —	ABSENT No environment-snapshot surface.	SPECCED **QBR Advisor coworker** is specced to "generate client-ready QBR narratives and security/posture summaries" — the closest conceptual match in the whole corpus. **0 of 40 tasks complete**; its seed migration never landed; the live `Coworker` model is an older, flatter shape than the spec assumes. `specs/004-ai-coworker-registry`	ABSENT —	ADJACENT **Value Ledger `GET /summary`** — the only reporting surface in the build. Automation ROI (time/cost saved by coworkers), not client business goals. `api/v1/value_ledger.py:93`

	ENGAGE — THE MEETING ITSELF (STEFANIE'S TAB ORDER)				
SERVICE MOMENT	Hardware/Software Asset Review prod-line	Employee Review prod-line	Technology Roadmap hybrid	Technology Budget brains	Additional Items unclass.
WHICH PULL					
ACTIONS (CUSTOMER)	Reviews asset/warranty status.	Reviews staffing notes.	Reviews 12–36mo roadmap; re-prioritizes live.	Reviews budget tied to roadmap. **CFO asks for a copy — the buy-in moment.**	*[?] unspecified content.*
PROCESSES, TOOLS, TECH					
THINKING & FEELING	*[?]*	*[?]*	*[?]*	*The buy-in feeling — CFO wants a copy.*	*[?]*
STRESSORS	*[?]*	*[?]*	*[?]*	*[?]*	*[?]*
OPPORTUNITIES (BUILD IMPLICATION)	Production-line — safe to auto-populate from pulls 1 & 3.	Production-line — safe to auto-populate from pull 3.	Re-prioritization logic could be assisted, but the priority call itself is judgment, set in Prep-Judge and re-tested live here.	Protect as brains-work — this is the clearest single ROI-demonstration moment to the client's own buyer (the CFO).	*[?] content unspecified.*
STAFF ACTIONS — FRONTSTAGE —	**Vanessa**: presents asset review.	**Vanessa**: presents staffing notes.	**Vanessa**: presents roadmap, re-prioritizes live.	**Vanessa**: presents budget, secures buy-in.	**Vanessa**: covers additional items.
STAFF ACTIONS — BACKSTAGE —					
SUPPORT PROCESSES					
ACP TODAY (IN THE BUILD)	ABSENT No asset inventory.	ABSENT No staffing surface.	ABSENT No roadmap object.	ADJACENT Value Ledger holds `cost_saved` + encrypted hourly rate — the only money in the build, but it is *our* ROI, not the client's IT budget. `value_ledger/models.py`	ABSENT —

	ENGAGE — THE MEETING ITSELF (STEFANIE'S TAB ORDER)		FOLLOW-UP	EXTEND
SERVICE MOMENT	Referrals & Testimonials brains	Next Meeting Planning prod-line	Post-Meeting Action Items + KPIs prod-line	steps 7–8, living document
WHICH PULL				
ACTIONS (CUSTOMER)	May give a referral or testimonial when asked.	Confirms next EBR date, live, closing the loop to Schedule.	Receives client-facing extract only.	Receives ongoing marketing/anniversary touches; benefits from continuity.
PROCESSES, TOOLS, TECH			Follow-up email, async, client-facing extract only.	Marketing touches, channel unspecified.
THINKING & FEELING	*[?]*	*[?]*	*[?]*	Arrives *guarded* next cycle if the MSP's gone quiet in between.
STRESSORS	*[?]*	"Client said no" taken at face value — a terrible reason per Stefanie.	**No follow-through** — action items never become tickets → "quickly cause your customer to lose faith." Stefanie's #1 killer.	Client refuses EBRs entirely (some do); switching-cost risk if Vanessa leaves.
OPPORTUNITIES (BUILD IMPLICATION)	A coworker could surface/remind the ask opportunity but shouldn't script the relationship-dependent moment itself.	Preserve the "book next meeting live" moment in any redesigned close — highest ROI-per-effort step, easiest to silently break.	⚙ **Action-item → PSA ticket**, already an explicit instruction — clearest win in the template. Build two real artifacts (working doc + extract), not one doc with a toggle.	Price/name the living document's switching-cost value. Any TAM math must respect the tiered/opt-out population.
STAFF ACTIONS — FRONTSTAGE —	**Vanessa**: asks for referral.	**Vanessa**: confirms next date live.	**Vanessa**: transcribes, reissues roadmap+budget, converts tickets, chases referral.	Ongoing marketing to existing base.
STAFF ACTIONS — BACKSTAGE —			*[?] unclear who enters tickets — Vanessa or a delegate.*	Institutional memory in the living document; cadence re-tiered per client (A/B/C; some refuse).
SUPPORT PROCESSES			PSA/Halo ticket queue.	The living/working document as an institutional-memory system — "a switching-cost surface nobody in canon has priced."
ACP TODAY (IN THE BUILD)	ABSENT —	ADJACENT ScheduleService again — the mechanism exists, unwired.	ADJACENT **Ticket creation is absent.** The nearest thing is a narrative idea about writing queue *re*-assignments back to ConnectWise — routing, not creating. Value Ledger's CSV/PDF export job is the only client-ready-artifact mechanism that actually ships. `value_ledger/templates/report.html`	ABSENT No cross-year living document. ("Institutional memory" appears in two Intent Specs — as a *changelog header for the spec document itself*. Different thing.)

What acp-core already has, checked against this map

Read at 2026-07-29 from `acp-core-main` — the Intent Specs, the five `specs/` folders, and `apps/backend`, distinguishing what is in code from what is only in prose. Bottom row of the grid above carries the per-step verdict. The short version: **of the 26 steps in this map, 0 are built, 1 is specced, 8 are adjacent, and 17 are absent.** That is not a criticism of the build — acp-core was scoped to autonomous runbook execution, not to account management. It means this map describes work the product does not currently touch.

The one real match — and it's a mechanism, not a feature

The **Value Ledger** is fully built (~2,100 lines across service, models, baseline resolver, rollup, and Temporal export workflows). Every entry carries a `baseline_source` tag resolved *once at write time* and permanently embedded, never re-resolved: `tenant_configured` → `platform_default` → `coworker_reported` → unknown. A model self-reporting zero saved minutes is rejected outright as a probable bug.

That is **provenance tiering already shipped in production code** — tenant-verified beats platform-estimated beats model-self-reported beats unknown-and-flagged. It is the same discipline this whole map is built on: never let a number be quoted without saying where it came from. The build already has the mechanism the EBR needs. What it doesn't have is any of the eight subjects the EBR needs it pointed at. **The gap is subject matter, not mechanism.**

The Ledger also emits a real client-ready artifact: `POST /exports` produces CSV or a Jinja/WeasyPrint-rendered PDF (sync under 1,000 entries, async Temporal workflow above it). So the build can already generate a document and hand it to a human — the pipe exists.

Named integrations in the build

INTEGRATION	STATUS	WHERE IT ACTUALLY LIVES
MCP server / instance framework	**BUILT**	A generic bring-your-own-connector layer — credential or OAuth, per-tool allowlist, HITL flag, tenant-scoped. **Ships with zero named connectors.** The migration defines the `mcp_servers/mcp_credentials` tables and no vendor rows. `specs/004-mcp-configs · app/mcp/`
ConnectWise Manage (PSA)	narrative only	Explicitly marked **"pending spec — not started"**. `Intent Specs/FSN-82692:437`
Autotask · HaloPSA (PSA)	named, not detailed	Listed as future PSA types in a non-goal. `Intent Specs/FSN-82692:272`
N-central / RMM	out of scope	"Not required… introduced post-onboarding." `Intent Specs/FSN-82692:269`
Jira · GitHub · Jamf	test fixtures	Arbitrary strings exercising the generic MCP model in unit tests and demo seed data. Not connectors. `tests/agents/test_expert_api.py · scripts/seed_experts.py`

No PSA, RMM, M365/IdP, backup, GRC, or license-management system is wired into this codebase. Every one of the eight prep pulls would need a connector registered from scratch. That is the honest size of the integration problem — and it is the same list, eight times over, that the Head of Technical Overview is currently doing by hand.

Three things worth knowing before Will and I decide

WHAT WE FOUND	WHY IT MATTERS TO THE BUILD DECISION
"QBR Advisor" is already a named coworker in the spec — one of five v1 platform defaults, scoped to "generate client-ready QBR narratives and security/posture summaries."	Someone already had this instinct and wrote it down. But **0 of 40 tasks are complete**, the seed migration never landed, and the live `Coworker` ORM model is a flatter, older shape with none of the `guardrails/trigger/autonomy_level` fields the registry spec assumes. So the name exists and the thing does not — which makes it easy to over-read as coverage in a status conversation. `specs/004-ai-coworker-registry/{spec.md:57, tasks.md:81}`
Scheduling is built but unwired. `ScheduleService` wraps Temporal Schedules, fires a coworker on a cron or interval, and has passing unit tests.	It is called from no API route and no coworker surface exposes a schedule field. The recurring-cadence mechanism an EBR needs is sitting there finished with nothing plugged into it. Cheapest real path from this map to shipped code probably runs through here. `app/runs/temporal/schedules.py · api/v1/coworkers.py`
The build's "onboarding" is the opposite direction from the map's.	FSN-82692 onboards *the MSP into ACP*. The Enter stage of this map is the MSP onboarding *its own client* — the step that sets the precedent for whether EBRs happen at all. Same word, inverted actor. Worth catching before anyone treats one as covering the other.

Absent, in full

Warranty/EOL · license-utilization / dormant seats · AD-IdP-vs-HR reconcile · ticket volume + SLA · patch compliance · backup job success and restore-test verification · phish-sim and awareness training · tenant posture vs. baseline · compliance checklist vs. a named standard · the Judge step · all 15 meeting tabs · PSA ticket *creation* · dual internal/client-facing artifact separation · the cross-year living document. Fourteen categories, no code and (except QBR Advisor) no spec.

Three findings that change what you should build

1 — The Risk Register's second job, and why Prep-Judge / Risk & Security Review are careful (unvalidated, kept in full)
Stefanie states the register's purpose three ways, none of them "prove value": shift the risk onto the customer, survive the insurer, stay out of court. The John anecdote (documented warning removed blame/emotion when a client got hit) is the mechanism appearing to work. **Downgraded 2026-07-28 after a legal scout pass:** two named cases exist, not three, neither 2025 (Mastagni Holstedt v. LanTech; Boardman Molded Products v. Involta); no case law, US or UK, has tested a documented refusal as a defense; no carrier treats it as mitigating — underwriters assess deployed controls, not paperwork; the defense gets pierced by the seller's own marketing (Involta's "one-stop shop" copy was used against it); the jurisdiction is inverted (advice is UK-sourced, litigation is US); and the mechanism runs both ways — machine-generated risk detection expands the record of what the MSP *knew*. **Volume of detection is volume of exposure.** Status: unvalidated, not a finding, not to appear client-facing before counsel.

2 — The 4–5h baseline's own credibility problem
Vanessa's 4–5h/EBR was originally read as independently corroborated by Stefanie's own figure. It is not — the source's own timeline shows Vanessa adopted this framework *after* Craig brought it home from Stefanie's training. Same process, measured twice at two points in its transmission — an echo, not validation. Wider vendor-blog figures (3–8h, "4–6 hours") have no primary source anywhere. Three unsourced numbers agreeing is not triangulation. Any reduction claim must be phrased as a ratio against this one-conversational-self-report baseline, never as an absolute — the 15-min-edit assumption in canon is really a claim of ~16–20× reduction on a job-shop task with a human-judgment core.

3 — The dual-surface artifact is structural, not a nice-to-have
Two living documents: Risk Register + Technology Roadmap/Budget feed The Working Document (internal notes Vanessa keeps) which produces The Client-Facing Extract (follow-up email only). Any product that renders *one* document — even with a redaction toggle — breaks this structurally. It needs to be two real artifacts. The 2026-07-28 Tier-1 prototype did exactly this wrong thing; council caught it (`prototypes/acp-qbr-ebr/council-pipeline-fast.md`).

Job shop vs. production line — every step, why it's sorted that way, and what good looks like

The distinction, in one paragraph. A **production line** step produces the same output every time from the same inputs — no judgment in the middle. Anyone (or anything) following the procedure gets the same answer, so the only questions are speed and cost. A **job shop** step produces a different output every time because the answer depends on *this* client, *this* quarter, *this* room — the value is the judgment itself, so making it faster by removing the judgment destroys the thing you were selling. (The framing is David Maister's leverage-ratio distinction from professional-services economics, applied to the EBR in `qbr-ebr.context.md` §6.)

Stefanie's two rules are the whole spec. On the production-line half: **"assemble it — and never show it."** On the job-shop half: **"this is the whole meeting."** Everything below is those two rules applied step by step.

Why this matters for what we build: automating the production line is a cost-and-speed play with a known-good answer. Touching the job shop is a different bet entirely — the tool can hold the structure and surface the prompt, but the moment it writes the judgment, the operator's own competence is what got automated, and the failure is silent. The one failure mode named more than any other in the source is an operator who *shows the production-line output instead of doing the job-shop translation* — so a product that makes the production-line output prettier and easier to display makes the most-cited failure mode *easier to commit*. That's the sharpest design tension on this page.

STEP	CLASS	WHY IT'S SORTED THAT WAY	WHAT GOOD LOOKS LIKE	IF WE BUILD IT WRONG
PREP — ASSEMBLE (THE 2–3H NOBODY SEES)				
Pulls 1–6 warranty/EOL · license use · user reconcile · ticket+SLA · patch+backup · phish-sim/training	prod-line	Each is a query against a system of record with one correct answer. No client-specific interpretation is needed to *produce the number*. Six different systems, one mechanical operation each.	All six land assembled, current, and complete before anyone opens their mouth — and a human never touches a keyboard to get them. Pull 5 reports backup *restore-tests* separately from backup-job-success; those are two different claims and only the weaker one is reliably in source.	Today's failure is not wrongness, it's *cost*: it's done by hand, late, by the one person who should be on the paying service desk. Build it so it's still manual-ish and you've kept the whole cost center.
Pull 7 tenant config vs. baseline	hybrid	The *diff* is mechanical. The *baseline* it's diffed against is the MSP's own authored standard — an opinion about what good security looks like for this kind of client, not a universal fact.	The comparison runs automatically; the baseline stays editable, attributable, and visibly the MSP's own. The MSP can see and defend which standard it held the client to.	Ship a vendor-defined baseline as if it were objective and you've quietly taken over the MSP's security opinion — and made every "you're non-compliant" claim our claim, not theirs.
Pull 8 compliance checklist	hybrid	Checklist status is pullable *if* the standard is codified. Which standard applies to this client, and whether they're "compliant enough," is judgment — and the source never names which standards apply per client.	Items auto-populate against a named, selected standard; the "are we OK" narrative stays a human sentence, written by someone who knows what the client is exposed to.	Auto-generate a compliance verdict and we've made a legal-adjacent assertion on the MSP's behalf with no counsel in the loop.
PREP — JUDGE (THE STEP WITH NO TOOL AT ALL)				
Translate gaps into business consequence · set roadmap priority	brains	This is the actual constraint in the whole process. Turning "17 machines are out of warranty" into "here's what that costs you if it fails in November" requires knowing this client's business. Source is explicit: **"you can't prepare 20 minutes before you walk in."** It is also the one step in the entire map with **no supporting tool named**.	The operator walks in with the translation already written and believes it. The tool's job is to hold the structure, put the right prompt in front of them, and refuse to let an unverified claim get spoken — *not* to write the sentence.	Two ways to fail. Draft the translation for them and their competence quietly erodes with no signal. Skip it and you've built a data-dump generator — which is the most-cited failure in the corpus.
THE MEETING — STEFANIE'S 15 TABS, IN HER ORDER				
Meeting Kick Off	unclass.	Never sorted in the original pass. Named honestly rather than guessed — setting the room's purpose is plausibly judgment, but the source doesn't say.	Unknown. Worth asking an operator directly before building anything here.	Guessing its class is how a build decision gets made on nothing.
Current Climate	prod-line	A snapshot of state. Fed almost entirely by pulls 1, 3 and 4 — the numbers are the content.	Auto-populated, zero prep, correct on arrival. The lowest-risk automation target on the whole map.	Little downside — but also little differentiation. Every competitor can do this.
Risk & Security Review	hybrid	The load-bearing moment. Entries come from pulls 5–8; the delivery is the business-consequence translation. Notably delivered by the client-delivery lead, **deliberately not a technician**.	Register entries arrive pre-populated and ranked; the room hears consequences, not findings; tech questions get parked rather than answered live. Any claim that hasn't been verified is gated *before* it's spoken.	This is where discomfort wins and raw data gets shown. Make the raw view the path of least resistance and we've built the failure mode a nicer front door. See finding 1 below on the register's second job — it is *not* validated.
Compliancy Review	hybrid	Same shape as pull 8: status is pullable, standing is a judgment call with legal adjacency.	Status shown against a named standard, verdict spoken by a human.	A generated verdict is an assertion we don't have standing to make.
Business Priorities & Goals	brains	There is no system of record for what a client is trying to do next year. The answer only exists in the conversation. Engagement here is also what makes the later referral ask land.	Real 12/24/36-month plans surface, in the client's words, and get captured verbatim into the living document. The tool listens and structures; it does not suggest goals.	Template this and it collapses into tech/stats recitation — and per source, the client sees no value and declines the next one.
Hardware/Software Asset Review	prod-line	Direct output of pulls 1 and 3. Inventory is inventory.	Fully auto-populated, warranty/EOL dates flagged by urgency, ready to feed the roadmap and the budget without retyping.	Leaving it manual burns technician hours for zero differentiation.
Employee Review	prod-line	Output of pull 3 — the joiners/leavers diff.	Auto-populated, with provisioning follow-ups falling out as action items rather than being noticed by a person.	Same as above: cost, not correctness.
Technology Roadmap	hybrid	The *entries* come from the gaps found in prep; the *order* is a judgment call that gets set in Prep-Judge and then re-tested live in the room against what the client just said their priorities were.	Entries flow in automatically; re-prioritizing live is fast and low-friction; the reordered roadmap is reissued after the meeting without retyping.	Auto-rank the roadmap and we've replaced the one judgment the client is actually paying for — and made live re-prioritization feel like fighting the tool.
Technology Budget	brains	Money tied to the roadmap, in front of the person who controls it. Source's clearest buy-in signal is the client's CFO asking for a copy — that's the ROI demonstration landing in real time.	Budget stays tied line-by-line to roadmap items so the client can see what each number buys. Protected as a live conversation, not a generated figure.	A generated budget nobody in the room owns is the fastest way to lose the one moment where value is visibly proven.
Additional Items	unclass.	Content unspecified in source. Not sorted.	Unknown.	Don't build to an empty slot.
Referrals & Testimonials	brains	Purely relationship-dependent. Whether to ask, and how, depends on how the last 50 minutes went — which is why it sits at the end of the agenda, not the start.	The tool *reminds* that the opportunity exists and tracks whether the ask happened. The operator decides and phrases it.	Script this and it reads as a script. A templated referral ask damages the relationship it's trying to monetize.
Next Meeting Planning	prod-line	Mechanically trivial — but per source it must happen **live, in the room, before anyone leaves**. Highest return per unit of effort in the whole process, and the easiest to silently break.	The date is committed verbally in the room and captured instantly. Cadence tier (A/B/C — not every client gets the same frequency, and some refuse entirely) is respected.	Replace the live verbal commit with a "we'll send an invite" and you may have broken the show-rate. Untested risk — flag it, don't assume it's safe.
Post-Meeting Action Items + KPIs	prod-line	Turning agreed actions into tracked tickets is pure mechanism — and the template already instructs it explicitly.	Every action item becomes a real PSA ticket with an owner, automatically. This is the single clearest win in the template.	**Stefanie's #1 killer.** No follow-through — action items that never became tickets — will "quickly cause your customer to lose faith." Get this wrong and the entire meeting was theatre.
AFTER — FOLLOW-UP AND THE LIVING DOCUMENT				
The follow-up extract	hybrid	Transcription and ticket conversion are mechanical. *Deciding what crosses the line to the client* is judgment — and the internal working document and the client-facing extract are two genuinely different artifacts, not one document with a filter.	Two real artifacts exist. The internal one keeps the operator's candid notes; the client one is composed, not filtered. See finding 3 below.	One document plus a redaction toggle breaks this structurally — an accidental un-toggle is a client-visible incident. Our own Tier-1 prototype made exactly this mistake.
The living document / institutional memory	brains	Its value is accumulated context — years of what this client said they wanted and what was decided. Nothing generates that; it's deposited over time.	It compounds. Continuity survives the operator leaving. Source calls this a switching-cost surface nobody in canon has priced.	Treat it as a meeting-notes archive and we've missed the one asset here that gets more valuable the longer a client stays.

Jargon index

TERM	PLAIN MEANING
EBR / QBR	Executive / Quarterly Business Review — the recurring meeting where an MSP sits down with a client's leadership and reviews the state of their technology, risks, and plans. The thing this whole map describes.
MSP	Managed Service Provider — the outsourced IT company. The customer of the product we'd be building.
RMM	Remote Monitoring & Management — the software an MSP uses to see and manage a client's computers from afar.
PSA	Professional Services Automation — the MSP's own ticketing/billing/ scheduling system. "Halo" is the specific PSA product named in source.
SLA	Service Level Agreement — the response/fix time the MSP contractually promised.
EOL	End-of-Life — the point where a manufacturer stops supporting and patching a product. After EOL, a known flaw can't be fixed at all.
Warranty (hardware)	Whether the manufacturer will still repair or replace the machine. Separate from EOL: a device can be in warranty but past EOL, or vice versa.
Tenant	A client's own dedicated slice of a shared cloud service (e.g. their own Microsoft 365 environment, separate from every other MSP client's).
Baseline	The MSP's own written standard for how a client environment should be configured. An opinion, not a universal fact — which is why pull 7 is hybrid.
Phish-sim	Simulated phishing — fake scam emails deliberately sent to a client's staff to see who clicks.
Patch compliance	What share of machines are up to date on security updates.
Restore test	Actually proving a backup can be recovered, as opposed to the backup job merely reporting success. Different claim; commonly skipped.
RTO / RPO	Recovery Time Objective / Recovery Point Objective — how long the client can afford to be down, and how much data they can afford to lose. Both are business decisions, which is why probing them is job-shop work.
IdP	Identity Provider — the system that manages who can log in (e.g. Active Directory, Okta).
GRC tool	Governance, Risk & Compliance software — tracks compliance requirements against a checklist.
Risk register	The running list of known risks in a client's environment, including ones the client was told about and declined to fix.
Cadence tier (A/B/C)	How often a given client gets an EBR. Not universal — some clients get them quarterly, some annually, some refuse entirely.
AM / AE	Account Manager / Account Executive — the relationship owner and the salesperson, respectively.
⚙ / ★ / ★★	⚙ = a machine can do this with no human judgment · ★ = needs a human's judgment call · ★★ = needs judgment from a specific, scarce person, not just anyone.