

Control Plane / Headless IT Tech in the AI Agent world

The Full Positioning Document

prototype of what it could look like

Why the MSP cares section

How Seth Godin would sell this

What we are building

Technical Overview

Agent types

Control plane responsibilities

Orchestration layer (first-party agents)

Management plane (vendor agents)

The ACP-First MSP: From Zero to Running Without Opening a Single Product (the ultimate aspiration)

Discovery

First Login: The ACP Dashboard

Customer Onboarding:

Customization: When Defaults Are Not Enough

Reporting and Value: Proving the Business

How does this work with N-zo and in the other products (Cove, UEM, Adlumin) in the near term?

Monitoring the Control Plane

Agent and Coworker Lifecycle

Execution Telemetry

Runbook Telemetry

Identity and Access Events

MCP and Tool Transaction Log

LLM Inference Telemetry

Policy and Governance Events

Value Ledger

It answers: what did the ACP actually do for this MSP, and what is that worth?

Monetization Metering (Billing Events)

Use Cases: Example with us and a Halo Expert

Set of 10 Use Cases we are sending to customers to rank:

Business Report for your customer

Alert Flood Suppression Coworker

Security Tabletop Exercise Simulator

Continuous Environment Health Coworker

MSP Employee Onboarding Coworker

Existing Customer New Employee Onboarding/Offboarding Coworker

New Customer Onboarding Coworker

Network Incident Resolution Coworker

MSP Benchmarking & Threat Hardening Coworker

Ticket Triage and Dispatch "Coworker" that:

MSP Acquisition Integration Coworker

Disaster Recovery Orchestration Coworker

Other use cases considered

Network Incident Resolution: Deterministic + Non-Deterministic in Action

Monetization: The ACP as a Revenue Architecture

Monetization: The ACP as a Revenue Architecture

- The Per-Device Model Is Broken
- How the Market Already Prices This
- Seven Revenue Layers (They Stack)
- How the Layers Compound
- How the MSP Charges Their Customer
- The (not named Marketplace) Vendor onboarding, access, and monetization process
 - Onboarding a Vendors' AI Expert
- Certification
 - Revenue Model
 - Versioning
 - The Marketplace
 - Ongoing Governance

The Full Positioning Document

[ACP_Positioning_Document.docx](#)

prototype of what it could look like

<https://n-able.atlassian.net/wiki/spaces/AI/pages/61995876354/ACP+prototype>

Thank you to @Elizabeth Connor

Why the MSP cares section

How Seth Godin would sell this

The MSP industry has a story it tells itself.

"We manage your IT." That is the story. And for twenty years, it worked. You charge per device, you hire technicians, you grow linearly. More customers, more people, more overhead. Repeat.

Here is the problem with that story: it is a commodity story. Every MSP tells it. The customer cannot tell the difference between you and the next one. So they shop on price. And you race to the bottom.

Now. What if you told a different story?

"We run 14 AI Coworkers across your environment. Last month they resolved 340 incidents, onboarded 12 employees, closed 2,100 vulnerabilities, and generated your quarterly business review. Here is the proof."

That is a story no one else is telling. That is a story the customer cannot compare to anything. That is a purple cow.

The Autonomous Control Plane is not a product. It is the infrastructure for a new story.

The old story sells time. The new story sells outcomes.

The old story says "trust us, we have good people." The new story says "here is a ledger of every action, every result, every dollar of value we created for you last month."

The MSPs who adopt this will not compete with the MSPs who do not. They will be playing a different game entirely. Their customers will not be able to go back to the old way, because the old way cannot show its receipts.

That is the shift. Not "we added AI." Nobody cares that you added AI. They care that you changed what you deliver, how you prove it, and what it is worth.

The ACP just makes that change possible at scale.

The interesting question is not whether the technology works. It does. The interesting question is whether you are ready to tell a different story.

What we are building

The ACP lets MSPs build, run, and govern AI agents that automate their operations, and we meter every execution.

It controls who they are, what they can do, what tools they can touch, and what happens when they run. Every action is logged, every permission is scoped, and every execution is metered.

MSPs use it to build AI Coworkers that automate real operational work: onboarding, triage, patching, incident response, reporting. They assemble Coworkers from N-able Experts, vendor Experts, their own runbooks, and their own tool connections.

The ACP is an independent product line, and may be purchased without other N-able products. However, it is our intent that our AI Experts and runbooks are so high value that they cant help but purchase them.

Technical Overview


The N-able control plane is the governing layer for AI Agents operating within the N-able product ecosystem. It serves three distinct but related functions: it is a full orchestration layer for N-able's own first-party AI agents (AI Experts), managing their execution path, resource allocation, and runtime policy; and it is a management plane for vendor-supplied agents and skills, governing their lifecycle, certification, and the conditions under which they operate within a customer environment and for Customer created AI Agents called AI Coworkers.

The control plane, and N-able generally, **does not build or host its own AI models** for this purpose. All inference is delegated to external LLMs, either provided through N-able's own processing infrastructure (aka Bedrock) or supplied directly by the customer.


The thing which is being controlled in this case are AI Agents, the Skills those AI Agents call, and the tooling which the skills utilize. These are at once:

- Non-Human Identities (**NHI**) as well as
- **Software** and each AI Agent is in itself
- a **mini-control plane**

We can imagine it looks a little like this:




with the AI Coworker detail view:



So the coworker ends up acting as a mini control plane on its own, managing workflows within the context of the MSP.

We can extend the concept so that “Once the coworkers works, you can export it as an MCP server so Claude can call it like a tool instead of making you rewrite the same prompt over and over. See :

https://app.therundown.ai/guides/build-a-custom-blog-writing-agent-with-no-code-langflow?utm_source=www.therundown.ai&utm_medium=newsletter&utm_campaign=zuckerberg-s-500m-ai-biology-swing&_bhlid=37e6fee453673027723409b34b64f4aaffd8aea7



([Experiment outline with Omni as a base](#))

AI Agent is the fundamental unit of the ecosystem. An agent combines an identity, a runbook, a set of tool access permissions, and a designated LLM into a deployable entity. Each deployed agent instance receives its own scoped identity there is no shared service principal across customers or deployments. This means two customers running the same Agent Expert are operating distinct principals with independent permission scopes.

Agent identity is modelled on the concept of an office worker with a defined role and access rights. Each deployment gets an identity that expresses what the agent is permitted to do, which systems it can reach, and under what conditions. Identity is enforced at runtime, not at configuration time.

[Runbooks describe what an agent will do \(can be as simple as a prompt or as complex as multiple SOPs\). Runbooks can be authored by N-able, by vendor partners, or by customers directly. They can be uploaded or generated. The control plane does not pre-screen or approve runbooks before use. This places responsibility on the runtime layer to catch out-of-policy behavior regardless of runbook provenance. This is a big 'unknown'. Runbooks may be as simple as a prompt, or as complex as containing multiple workflows. PLEASE SEE THE RUNBOOK WRITEUP \[HERE\]\(#\)](#)

MCP servers (and CLI and APIs and whatever mechanism comes next) are the mechanism by which agents access software systems. They form the tool access layer sitting between an agent and the services it needs to operate. Which MCP servers an agent may use, and with what scope, is governed by the control plane (particularly for vendor agents, where permitted servers and scopes are explicitly bounded).

MCP keys and credentials are managed by a dedicated key service, separate from the control plane itself. The key service is responsible for credential vaulting and issuance, and issues credentials scoped to a specific deployment rather than to an agent type. This means credential rotation and revocation operate at the deployment level.

LLMs are always external. N-able does not build or operate its own models. Customers may use N-able's processing infrastructure (which routes to a supported external model) or bring their own model and API key (BYOK/BYOM). In the bring-your-own case, the control plane hands off to the customer's chosen model and does not sit in the inference path. LLM selection is a per-agent configuration.

The script repository provides deterministic, non-AI tooling that agents can invoke for activities where predictable, auditable execution is required. This is distinct from MCP-based access and represents a lower-variance execution path for well-defined operations.

Agent types

AI Agent Experts are fully preconfigured AI agents released by N-able. They carry a complete end-to-end use case, a fixed runbook, defined tool access, and an N-able-managed identity. They require no customer configuration to deploy and operate entirely within the orchestration layer.

Approved Vendor Experts are AI agents built by third-party vendor partners and certified by N-able for use within the ecosystem. Certification involves N-able reviewing the agent's intended behavior, permitted MCP scopes, and compatibility with ecosystem policy. Once certified, a Vendor Expert is listed in the agent catalogue and available to customers. Vendor Experts are managed through the management plane rather than the orchestration layer. N-able governs their lifecycle, versioning, and policy compliance, but does not necessarily own their runtime.

AI Coworkers are user-assembled super-agents. A customer builds a Coworker by combining their own runbooks, their own MCP configurations, and any mix of Agent Experts and Approved Vendor Experts. The result is a custom workflow tailored to the customer's environment. Coworkers execute within the orchestration layer, but the vendor agents they embed may execute either on N-able infrastructure or on the vendor's own infrastructure, at the customer's discretion.

Control plane responsibilities

Orchestration layer (first-party agents)

For Agent Experts and Coworkers executing on N-able infrastructure, the control plane owns the full execution path. Its responsibilities include:

- Routing agent requests to the appropriate LLM endpoint, whether N-able-hosted or handed off to a customer-supplied model
- Enforcing agent identity and permission boundaries at runtime
- Brokering MCP server access by requesting appropriately scoped credentials
- Applying runtime policy against agent actions ***If it owns the*** runbook provenance
- Providing execution context and resource management for the agent's lifecycle

Management plane (vendor agents)

For Approved Vendor Experts, the control plane acts as a governance and certification authority rather than a runtime executor. Its responsibilities include:

- Operating the agent catalogue: onboarding, versioning, and certification of vendor agents
- Allowing the vendor to define permitted MCP server access and credential scopes for each vendor agent

- Establishing the policy conditions under which a vendor agent may be embedded into a AI Coworker

When a vendor agent executes on the customer's own infrastructure, the control plane governs ???(BIG QUESTION TO ANSWER HERE ... ? credentials ?, ?what it is permitted to request?, and ?how its outputs are validated? before they re-enter a Coworker's execution context. *The vendor's runtime is trusted within the bounds N-able has defined, but not unconditionally.*

Audit and observability provides logging, tracing, and policy check records across all agent activity. Because runbooks are not pre-approved, the audit layer is the primary mechanism for detecting and investigating out-of-policy behavior after the fact.

[There is an opportunity for a Deterministic Substitution Engine](#) which intercepts requests and transforms LLMs calls into script or API calls (pulling from the repo of available deterministic outcomes)

The ACP-First MSP: From Zero to Running Without Opening a Single Product (the ultimate aspiration)

If the control plane works, a new MSP should never need to learn UEM. They should never need to learn Cove. They should never need to learn Adlumin. Those products become infrastructure, engines running underneath, while the ACP is the only interface the MSP touches.

The MSP does not buy a collection of tools and figure out how to stitch them together. They buy a business-in-a-box that happens to run on N-able's product lines underneath.

Discovery

An MSP owner hears about the ACP the same way they hear about everything: a peer mentions it, they see it at a conference, or they find it searching for "how to start an MSP" or "automate MSP operations."

They land on a page that says: **Build AI Coworkers that run your MSP operations across every customer, every tool, every shift.** They do not see a pitch for endpoint management software. They do not see a pitch for backup software. They see a pitch for running an MSP.

They sign up for a trial.

First Login: The ACP Dashboard

The MSP logs into the ACP. There is no redirect to UEM. No redirect to Cove. No redirect to Adlumin. The ACP is the front door and the only door they need.

What they see:

- **A setup N-zo** that asks three questions: How many customers do you have today? What services do you plan to offer (endpoint management, backup, security, lists all existing Vendor Experts)? Do you have existing tools or are you starting fresh?
- **A catalogue of pre-built AI Experts & Coworkers** organized by MSP function: onboarding, monitoring, patching, backup, security, triage, reporting
- **A getting-started N-zo experience** that walks them through their first deployment in under an hour

Based on their answers, we provision the underlying product infrastructure automatically. If they selected endpoint management, UEM is provisioned behind the scenes. If they selected backup, Cove is provisioned. Security, Adlumin. The MSP never sees these provisioning steps. They see: "Your environment is ready."

Customer Onboarding:

In the ACP, they click **Activate** on the **New Customer Onboarding Coworker**. It is a pre-built, N-able-authored Coworker that ships with the platform. It comes with a default runbook that covers the standard onboarding flow.

The MSP provides the basics: customer name, primary contact, number of endpoints, service tier. The Coworker does the rest.

Behind the scenes, the Coworker delegates to default Experts:

- The **Endpoint Expert** provisions the customer in UEM, generates the agent installer, and sets default monitoring policies. The MSP never opens UEM.
- The **Backup Expert** creates the customer in Cove, configures default backup schedules based on the service tier, and validates the first backup completes. The MSP never opens Cove.
- The **Security Expert** enrolls the customer in Adlumin, deploys EDR, and applies the baseline security policy. The MSP never opens Adlumin.
- The **Documentation Expert** creates the customer record, logs every action taken, and generates a customer-ready onboarding summary.

The MSP sees a single execution log in the ACP: "Customer onboarded. 47 actions completed. 0 failures. Onboarding summary ready for delivery." They click, review the summary, and send it to their customer.

They just onboarded a customer across three product lines without opening any of them.

Customization: When Defaults Are Not Enough

Eventually, the MSP wants to customize. Their biggest customer has a specific onboarding policy. Their vertical requires a particular compliance report. They want alerts from a third-party tool correlated alongside N-able alerts.

They open the Coworker Builder in the ACP. They clone the default Onboarding Coworker. They modify the runbook to add the custom steps. They add a Vendor Expert from the marketplace (the Halo Ticket Expert, the SentinelOne Expert, whatever they need). They test it in sandbox mode. They publish it.

At no point do they open UEM, Cove, or Adlumin to make this customization. The ACP is the authoring environment. The products are the execution infrastructure.

If they write a custom script for a specific remediation, they write it in the ACP's script repository. It becomes available to any Coworker. If they build a custom runbook for a niche use case, it lives in the ACP and any Coworker can reference it.

Reporting and Value: Proving the Business

The MSP opens the Value Ledger in the ACP. It shows, across all customers, all Coworkers, all time periods:

- Executions run
- Incidents resolved without human intervention
- Employees onboarded and offboarded
- Vulnerabilities remediated
- Backup failures caught and fixed
- Manual equivalent hours saved
- Estimated cost avoidance

They pull a report for a single customer. It shows everything the automation did for that customer last month. They attach it to the QBR the Business Report Coworker already generated. They walk into the customer meeting with proof of value they assembled in five minutes.

How does this work with N-zo and in the other products (Cove, UEM, Adlumin) in the near term?

for simplicity, we will just call Cove, UEM, Adlumin “apps”

ACP → Apps (Publish)

- ACP is the authoring environment for Coworkers and Experts
- Publish targets include N-zo in UEM, Cove, and Adlumin
- Each Coworker is flagged as autonomous-only, interactive-compatible, or both
- N-zo invokes interactive-compatible Coworkers using a delegation model that respects both the Coworker's NHI and the calling user's permissions
- Autonomous Coworkers are available as triggers or background actions, not conversational experiences


Apps → ACP (Reverse Publish)

- Scripts, skills, and runbooks created in any product are available in the ACP
- Each carries provenance metadata: authoring product, author type (product team vs. customer), product context assumptions, cross-product validation status
- Assets are versioned independently; the ACP pins to a version and surfaces updates without forcing them
- Permission-gated: the logged-in user must have rights in both the source product and the ACP

N-zo's Role

- N-zo remains the in-app assistant with product-scoped context as its primary function
- Published Coworkers appear as a clearly separated capability layer within N-zo, visually and functionally distinct from native app assistance
- N-zo can trigger both interactive and autonomous Coworkers, with appropriate UX for each (conversational response vs. "kicked off, I will notify you when complete")

There is also a world in which the customers are ACP customers only and you could imagine N-zo as exposed through the ACP



Monitoring the Control Plane

The job of monitoring is to capture every agent action, every tool call, every credential issuance, and every customer outcome, categorize it, and make it usable for three distinct audiences: operations (is it working), finance (what do we charge), and compliance (can we prove what happened).

Agent and Coworker Lifecycle

Every entity that exists inside the ACP has a lifecycle. Creation, configuration, deployment, modification, retirement. All of it gets tracked.

- **AI Coworkers created** — who created it, when, which tenant, which Experts and runbooks were embedded, which MCP servers were connected
- **AI Coworkers modified** — what changed, who changed it, previous version vs. current version, timestamp
- **AI Coworkers deleted or retired** — when, by whom, reason if provided, whether active executions were in flight
- **Agent Experts activated** — which Expert, which tenant, deployment-scoped identity issued, LLM configuration at activation
- **Vendor Experts onboarded** — certification date, permitted MCP scopes, policy constraints, version
- **Vendor Expert version changes** — old version, new version, re-certification status, rollout scope

- **Coworker export events** — when a Coworker is exported as an MCP server (per the export model described elsewhere), who exported it, what was exposed

Execution Telemetry

Every workflow run, every Expert delegation, every step inside a Coworker.

- **Coworker executions** — trigger event, start time, end time, duration, completion status (success / partial / failure / timeout)
- **Runbook invocations** — which runbook, version, provenance (N-able authored, vendor authored, customer authored), whether the runbook was modified since last run
- **Expert delegations within a Coworker** — which Experts were called, in what order, parallel vs. sequential, per-Expert start/end/duration/status
- **Step-level trace** — every discrete action inside an execution: tool call, LLM inference, script execution, approval gate, conditional branch taken, data transformation
- **Trigger source** — what initiated the execution (scheduled, event-driven, manual, API call, upstream Coworker)
- **Retry and fallback events** — when a step fails and retries, how many attempts, whether a fallback path was taken, final disposition
- **Queue and wait times** — time between trigger and execution start, time spent waiting for approvals, time spent waiting for external system responses

Runbook Telemetry

Runbooks are the behavioral core of every agent.

- **Runbook inventory** — all runbooks in the system, version, author, provenance, last modified date, associated Coworkers
- **Runbook runs** — count per runbook, per tenant, per time period
- **Runbook drift** — when a runbook is modified between runs, delta between versions
- **?Runbook complexity metrics** — number of steps, number of conditional branches, number of tool calls, number of approval gates, estimated token cost per run
- **Runbook failure rate** — per runbook, per version, over time, with root cause tagging (tool failure, LLM error, policy block, timeout, bad input)
- **Runbook coverage** — which operational areas have runbooks, which are still handled manually, gap analysis over time

- **Out-of-policy events by runbook** — when a runbook execution triggers a policy violation, tagged to the specific runbook and version

Identity and Access Events

Every agent is a Non-Human Identity.

- **Identity issuance** — when a deployment-scoped identity is created, for which agent, in which tenant, with what permissions
- **Identity modification** — scope changes, permission additions or removals, who authorized the change
- **Identity revocation** — when and why an identity was revoked, whether active sessions were terminated
- **Credential requests to the key service** — which agent requested credentials, for which MCP server, scope requested vs. scope granted
- **Credential rotation events** — automated rotations, forced rotations, rotation failures
- **Credential revocation** — when, why, which deployment, whether dependent executions were affected
- **Scope violations** — when an agent attempts an action outside its permitted scope, whether it was blocked, what it attempted

MCP and Tool Transaction Log

Every interaction between an agent and an external system flows through the tool access layer. Every one gets logged.

- **MCP server calls** — which server, which agent, which method, request payload summary (not full payload for privacy), response status, latency
- **API calls** — endpoint, agent, method, status code, latency, payload size
- **Script executions** — which script from the repository, which agent invoked it, input parameters, execution result, runtime duration
- **Deterministic Substitution Engine intercepts** — when the engine intercepts a request that would have gone to an LLM and routes it to a script instead: original intent, substituted script, result, tokens saved
- **Tool failures** — which tool, error type, whether the failure was transient or persistent, retry outcome
- **Tool access denied** — when an agent requests a tool it is not permitted to use, logged with full context

- **Tool latency and availability** — per MCP server, per API endpoint, over time, for SLA tracking and capacity planning
- **Transaction volume by tool** — counts per MCP server, per API, per script, per tenant, per time period

This is the **metering backbone for Layer 4 [monetization](#)** (MCP and Tool Transactions)

LLM Inference Telemetry

- **Inference requests** — which agent, which model, prompt token count, completion token count, total tokens, latency
- **Model selection** — which LLM was used (N-able Bedrock vs. BYOK/BYOM), per agent, per execution
- **Inference cost** — calculated cost per call based on model pricing, aggregated per execution, per Coworker, per tenant, per billing period
- **Inference failures** — model errors, timeouts, rate limit hits, fallback behavior
- **Deterministic substitution rate** — percentage of potential LLM calls intercepted by the Substitution Engine, tokens saved, cost avoided
- **Model quality signals** — when available: confidence scores, hallucination flags, policy check failures attributable to model output
- **BYOK/BYOM routing events** — when inference is handed off to a customer-supplied model, logged for audit even though N-able does not sit in the inference path

This is the **metering backbone for Layer 5 monetization**

Policy and Governance Events

- **Policy violations detected** — what was attempted, which policy was violated, whether the action was blocked or flagged for review
- **Escalations to human** — when and why an agent escalated (low confidence, policy requirement, separation of duties), who received the escalation, resolution

Value Ledger

It answers: what did the ACP actually do for this MSP, and what is that worth?

- **Manual equivalent hours saved** — per Coworker execution, estimated time a technician would have spent performing the same work manually, aggregated per tenant, per billing period. May be taken from the Runbook estimated time for task

- **Incidents resolved without human intervention** — count, severity distribution, mean time to resolution vs. historical manual baseline
- **For Experts where we know the intended outcome:**
 - **Tickets triaged and routed** — count, accuracy rate (did the triage match the eventual resolution category), time saved per ticket - if we have the ticket data
 - **Remediations executed** — patches applied, vulnerabilities closed, configurations corrected, drift remediated, with before/after posture scores
 - **Employees onboarded/offboarded** (for — count, time from trigger to completion, systems provisioned or deprovisioned
 - **Customers onboarded** — count, time from trigger to completion, controls verified
 - **QBRs and reports generated** — count, data sources integrated, customer delivery confirmation
 - **Alert noise reduced** — alerts ingested vs. alerts surfaced to humans, suppression ratio, false positive rate
 - **Cost avoidance** — estimated cost of incidents prevented, breaches avoided (based on industry benchmarks), SLA penalties avoided
 - **Revenue influenced** — upsell opportunities surfaced by Coworkers (e.g., QBR Coworker identifies posture gaps that lead to new service sales), tracked if the MSP closes on them
- **LLM cost savings from deterministic substitution** — tokens that would have been consumed, actual tokens consumed, delta

The Value Ledger is customer-facing. It is the artifact the MSP shows their customer to justify the price. It is also the artifact N-able shows the MSP to justify the platform fee.


Monetization Metering ([Billing Events](#))

This is the subset of telemetry that directly feeds invoicing.

- **Execution count** — per Coworker, per tenant, per billing period (Layer 2)
- **Expert activations** — which Experts are active, per tenant, activation and deactivation dates (Layer 3)
- **MCP and tool transaction count** — per server, per API, per script, per tenant, per billing period (Layer 4)
- **LLM token consumption** — per model, per tenant, per billing period, with cost calculation (Layer 5)
- **Outcome events** — per resolved incident, per onboarded employee, per generated report, per billing period (Layer 6, future state)

- **Tier and entitlement tracking** — which tier the tenant is on, included allotments, overage thresholds, current consumption vs. entitlement
- **Billing anomaly flags** — sudden consumption spikes, usage patterns inconsistent with tenant size, potential misconfigurations driving unnecessary cost

Use Cases: Example with us and a Halo Expert



A managed service provider (MSP) uses Halo as their ticketing system. When a new employee joins one of their customers, someone raises a ticket in Halo to kick off the onboarding process.

The MSP has built an Onboarding AI Coworker in N-able. They wrote a runbook for it that says: when an onboarding ticket comes in, provision an identity for the new user and allocate them an asset. They've embedded the N-able Identity Expert and Asset Expert into the Coworker to handle those two jobs.

Here's what happens end to end:

Halo's Ticket Expert — a vendor agent certified in the N-able ecosystem ingests the ticket and passes it to the control plane. The management plane validates the payload and checks it

against policy before anything proceeds.

The control plane then triggers the Onboarding Coworker. It reads its runbook, sees it needs to do two things, and delegates to the Identity Expert and Asset Expert in parallel. Before either Expert does anything, the orchestration layer issues each one its own scoped identity and fetches the right MCP credentials from the key service.

The Identity Expert provisions the new user's identity in N-able. The Asset Expert allocates a device or asset to them.

When both complete, the Coworker collects the results and confirms the runbook is satisfied. The full execution trace is written to audit.

Finally, the result is passed back to the Halo Ticket Expert, which updates the ticket to show onboarding is complete.

Set of 10 Use Cases we are sending to customers to rank:

(workign document [Top 10 ACP Use Cases.docx](#))

[Business Report for your customer 1](#)

[Alert Flood Suppression Coworker 2](#)

[Security Tabletop Exercise Simulator 3](#)

[MSP Employee Onboarding Coworker 3](#)

[Existing Customer New Employee Onboarding/Offboarding Coworker 4](#)

[New Customer Onboarding Coworker 4](#)

[Network Incident Resolution Coworker 5](#)

[MSP Benchmarking & Threat Hardening Coworker 5](#)

[Ticket Triage and Dispatch “Coworker” that: 6](#)

[MSP Acquisition Integration Coworker 7](#)

[Disaster Recovery Orchestration Coworker 7](#)

Business Report for your customer

Integrates: PSA, RMM/UEM, Security (MDR), Backup, Reporting , etc.

SUMMARY FOR SURVEY:

What do these agents do?

Generate a customer-ready report with charts, summaries, and recommended next steps. It shows progress, risk reduction, and remaining gaps in clear business terms, using data pulled from across MSP tools to highlight trends in patching, vulnerabilities, incidents, backups, and ticket volume.

What's in it for you?

Turns QBR prep from days of manual work into minutes. It helps shift customer conversations from ticket noise to business outcomes, clearly proving the MSP's impact while naturally uncovering growth opportunities and building trust through professionalism and transparency.

Explanation:

- Pulls operational, security, and service data from across all MSP vendor tools
- Normalizes metrics across vendors so results are consistent and comparable
- Analyzes trends in patching, vulnerability exposure, incidents, backups, and ticket volume
- Translates raw metrics into clear business-level insights and narratives
- Highlights risk reduction, improvements made, and gaps that remain
- Generates a customer-ready report with charts, summaries, and recommended next steps

What's in it for you:

- QBRs take minutes instead of days of manual data gathering
- Conversations shift from tickets and noise to outcomes and value delivered
- The MSP clearly proves its impact on security, uptime, and risk reduction
- Upsell and expansion opportunities are surfaced naturally through data

- Customers see professionalism and transparency, increasing trust and retention

Alert Flood Suppression Coworker

Integrates: MDR tools, Backup tools, AlertOps, PSA tools, RMM alerts

SUMMARY FOR SURVEY

Cuts alert noise by correlating duplicate events, suppressing known benign or cascading alerts, and grouping related issues into a single problem record. It escalates to a human when patterns change or confidence drops, helping the team stay focused on what actually needs attention.

What's in it for you? **It reduces alert noise, lowers burnout, and improves response time.**

- Correlates duplicate alerts across tools and devices
- Suppresses known-benign or cascading alerts under policy
- Groups alerts into a single problem record
- Escalates when patterns shift or confidence drops

What's in it for you: Less noise, lower burnout, better response time

Incident Response Root Cause Analysis

Integrates: Adlumin XDR/MDR, RMM/UEM, Identity Provider, Backup (Cove), PSA, Documentation (Hudu)

SUMMARY

Generates a customer-ready Root Cause Analysis report with a clear timeline, key findings, actions taken, and recommended hardening steps. It automatically captures the full incident as it happens, connects every detection and response into a single chain of evidence, and traces the root cause across security, identity, endpoint, and network data. It also identifies contributing factors and updates documentation with lessons learned and improved SOPs.

What's in it for you? Builds the RCA as the incident happens, so the report is ready in hours instead of days. Every action, decision, and tool call is automatically logged and traced, giving the MSP a professional, defensible report without the usual manual digging. The result is faster delivery, stronger customer trust, and a natural conversation about how to prevent the issue from happening again.

Explanation:

Captures the full incident timeline automatically as events unfold across tools

Maps every detection, decision, approval, tool call, and remediation action to a single chain of evidence

Correlates the root cause by tracing backward from impact to origin across security, identity, endpoint, and network data

Identifies contributing factors: missed patches, policy drift, access gaps, configuration weaknesses

Generates a customer-ready Root Cause Analysis report with timeline, findings, actions taken, and recommended hardening steps

Updates documentation systems with lessons learned and revised SOPs

What's in It for you:

Today, writing an RCA takes hours of digging across tools, Slack threads, and ticket notes, and the result is inconsistent and often late. ACP builds the RCA as the incident happens, not after. Every action is already logged, every decision already traced, every tool call already timestamped. The MSP delivers a professional, defensible report to the customer within hours instead of days. That report becomes a trust builder, an insurance artifact, and a natural upsell conversation: "Here's what happened, here's what we did, and here's what we recommend to prevent it."

Security Tabletop Exercise Simulator

Integrates: N-able MDR, UEM, Backup, Vendor Security Expert, N-able Vulnerability Expert, N-able Script Expert, Script Repository

SUMMARY

Simulates a real security incident using the MSP's actual tools, SOPs, and environment, with realistic alerts instead of hypothetical scenarios. It walks through detection, investigation, and response exactly as they would happen, evaluating decisions, approvals, and escalation paths along the way. The result is a clear record of what happened, where gaps exist, and where automation or controls should be improved.

What's in it for you? Lets teams rehearse real incidents before they happen, so gaps in SOPs, permissions, and tooling surface safely instead of during a breach. It builds muscle memory without risking customer environments, gives leadership clear proof of readiness, and helps the MSP strengthen customer trust by demonstrating real response maturity.

Explanation:

Simulates a real security incident using the MSP's actual tools, SOPs, and environment

Injects realistic alerts and signals instead of hypothetical scenarios

Uses AI reasoning to walk through detection, investigation, and response paths

Evaluates decision points, approvals, and escalation flows exactly as they would occur

Identifies when remediation actions would be deterministic versus require human approval

Documents actions taken, gaps found, and where automation or controls should improve

What's in It for the you:

Rehearse real incidents before they happen

Gaps in SOPs, permissions, and tooling surface safely instead of during a breach

Teams build muscle memory without risking customer environments

Leadership gains clear evidence of readiness and improvement areas

The MSP strengthens trust with customers by proving response maturity, not just claiming it

Continuous Environment Health Coworker

Integrates: RMM/UEM, Security (Adlumin / EDR), Backup (Cove), PSA

SUMMARY

Runs continuous health sweeps across every customer environment to catch issues early and keep systems aligned with policy. It checks patching, backups, security posture, agent health, and core system performance, automatically fixing routine drift and opening tickets only when

human judgment is needed. The result is a daily summary of what was found, what was fixed, and what still needs attention.

What's in it for you? Keeps your customer environments healthy without requiring technicians to start each day by checking dashboards across multiple tools. It automatically handles routine drift before anyone logs in, reducing manual effort and helping the MSP deliver more consistent operations at scale. The result is fewer missed issues, less repetitive work, and a more reliable service experience for customers.

Explanation:

Runs continuous health sweeps across every customer environment on a defined schedule
Checks patch compliance, agent health, backup success, security posture, disk/memory/CPU, and certificate expiry

Automatically remediates drift: restarts agents, retriggers backups, pushes missing patches, clears temp files

Opens tickets only for issues that require human judgment

Produces a daily operations summary: what was found, what was fixed, what needs attention

What's in It for the MSP:

Today a technician spends the first hour of every morning checking dashboards across tools. ACP eliminates that entirely. The environment is already healthy before anyone logs in. The MSP replaces a daily human ritual with a persistent, cross-product coworker that never misses, never forgets, and never calls in sick.

MSP Employee Onboarding Coworker

Integrates: HR system, Identity Provider, PSA, RMM/UEM, Security tools

SUMMARY

Automatically starts onboarding across all required systems when a new employee is added. It provisions identity, device, and application access based on role and policy, validates security posture before access is granted, and requires approval for any sensitive permissions. Every step is logged for traceability and compliance.

What's in it for you? Gets new hires productive on day one without manual coordination or missed steps. It reduces security risk by making access consistent and policy-driven, cuts repetitive setup work for the team, and gives leadership clear proof that onboarding was done correctly and safely.

Explanation:

When a new employee is added, this coworker automatically initiates onboarding across all required systems. It provisions identities, devices, and application access based on role and policy, enforces least-privilege permissions, and validates device and security posture before access is granted. Any elevated or sensitive access requires approval, and every step is logged for traceability and compliance.

What's in It for the you:

New hires are productive on day one without manual coordination or missed steps. Security risk is reduced because access is consistent and policy-driven, not handled ad hoc. The team spends less time on repetitive setup, senior techs aren't needed for every onboarding, and leadership has clear proof that access was granted correctly and safely.

Existing Customer New Employee Onboarding/Offboarding Coworker

Integrates: Customer HR System, Identity Provider (IdP), Backup, RMM/UEM, Security (EDR), M365/SaaS, PSA

SUMMARY

Automatically detects when a customer adds or offboards an employee and starts the right onboarding or offboarding workflow based on customer and MSP policy. It provisions or removes access with least-privilege defaults, validates device and security readiness, and requires approval for any sensitive permissions. Every step is confirmed and logged for audit, compliance, and consistency.

What's in it for you? Gets customer employees productive faster without manual back-and-forth or inconsistent handoffs. It reduces security risk by making access and device readiness consistent across customers, helps the MSP scale onboarding and offboarding more safely, and provides clear proof that every step was completed correctly and compliantly.

Explanation:

Detects when a customer adds / offboards a new employee via HR or service request
Resolves role, department, location, and risk profile using customer and MSP policy
Provisions identity, device access, and required SaaS applications with least-privilege defaults
Validates device readiness and security posture before granting access
Requires approval for elevated or sensitive permissions
Confirms completion and logs all actions for audit and compliance

What's in It for you:

New customer employees are productive on day one without manual back-and-forth
Security risk is reduced by enforcing consistent access and posture across customers
Onboarding becomes repeatable and scalable instead of ticket-driven and error-prone
Junior techs can handle onboarding safely without senior oversight
The MSP demonstrates operational maturity with clear evidence of secure, compliant onboarding

New Customer Onboarding Coworker

Integrates: PSA, RMM/UEM, Security (MDR), Backup, Documentation, Reporting

SUMMARY

Starts a standardized onboarding flow across every connected system when a new customer is signed. It sets up the customer in the PSA, deploys baseline configurations across core tools, applies standard policies, verifies critical controls, and keeps risky changes behind approvals. It finishes with a customer-ready summary showing what was deployed, what was verified, and what still needs human input.

What's in it for you? Gets new customers onboarded faster without relying on manual coordination across tools and teams. It makes onboarding more consistent, reduces the risk of missed controls or unsafe changes, and gives the MSP clear proof that standards were applied from.

Explanation:

When the MSP signs a new customer, this coworker kicks off a standardized onboarding flow

across every connected system. It creates the customer record and onboarding project in the PSA, then provisions the tenant and baseline configurations across RMM/UEM, security, and backup tools.

It automatically deploys agents, applies standard policies, sets monitoring thresholds, and verifies critical controls are active (backup success, endpoint protection present, alert routing working).

Risky changes (admin access, firewall modifications, exception policies, broad scripting) are gated behind approvals, and the coworker logs every step with timestamps for auditability.

It finishes by generating a customer-ready onboarding summary: what was deployed, what was verified, what still needs human input, and the proof that standards were met.

Network Incident Resolution Coworker

Integrates: N-able, Vendor Network Expert, MDR, Backup, Vulnerability Expert, N-able Script Expert, Script Execution Engine

Detects network issues and triggers the customer's AI Coworker using their SOPs, then gathers the topology, device state, and logs needed to diagnose the problem. It uses AI reasoning to choose the right remediation, then shifts to validated script execution as soon as the path is deterministic, avoiding generated commands. Every action and outcome is logged end to end for speed, safety, and auditability.

What's in it for you? Resolves network incidents faster without manual triage or constant tool switching. It reduces risk by executing approved, validated scripts instead of generated commands, lowers AI costs by stopping model use once execution is deterministic, and creates a clear audit trail of what changed, why, and with what result.

Explanation:

Detects network errors and triggers the customer's AI Coworker using their SOPs

Orchestrates vendor and N-able experts to gather topology, device state, and logs

Uses AI reasoning to diagnose the issue and select the correct remediation

Identifies when an action can be executed deterministically

Executes validated scripts directly without LLM-generated commands

Stops model usage once deterministic execution is possible

Logs execution results and outcomes end-to-end

What's in It for you:

Faster incident resolution without manual triage or tool switching

Lower risk by executing known-safe, approved scripts instead of generated commands

Reduced AI costs by eliminating unnecessary token usage during execution

Repeatable outcomes that don't depend on senior technician intervention

Clear audit trails that prove what changed, why, and how

MSP Benchmarking & Threat Hardening Coworker

Integrates: MDR, UEM, Vulnerability scanner, EDR/Security tools, Backup, PSA/Reporting

Pulls together patching, vulnerability, and security posture data across your customer base and normalizes it so comparisons across vendors are actually meaningful. It places you into a peer group based on MSP profile and benchmarks you on key operational metrics, then highlights where you're ahead, where you're behind, and which controls will improve results fastest. It finishes with customer-ready charts and messaging for QBRs.

What's in it for you? Shows exactly how your MSP compares to similar peers, making it easier to prove progress and explain priorities to customers. It turns benchmarking into a practical hardening tool by spotlighting the gaps that matter most and giving you a clear story about what you're improving next.

Explanation:

- Ingests your patching, vulnerability, and security posture across your full customer profile
- Normalizes data across vendors so “apples to apples” comparisons are actually possible
- Segments you into a “peer group” based on MSP size, customer mix, industry, endpoint count, and tool stack
- Benchmarks you against peers on daily operational metrics, for example: patch compliance timelines, critical vuln exposure window, EDR coverage, backup success rate, MFA/identity hygiene
- Flags where you're outperforming, where you're behind, and what specific controls would move the needle fastest
- Produces a customer-ready narrative and charts for QBRs, including “here's where you stand vs similar MSPs” and “here's what we're improving next”

Ticket Triage and Dispatch “Coworker” that:

Integrates: PSA/Helpdesk, RMM/UEM, Security (MDR), Backup, Identity

SUMMARY

Ingests new tickets from any PSA or helpdesk and adds live endpoint, user, and tenant context so triage starts with the right information. It uses bounded AI reasoning to classify severity, category, and likely root cause, then runs safe diagnostics automatically while stopping before any disruptive remediation. When confidence is low or policy requires separation of duties, it escalates to a human technician.

What's in it for you? Routes tickets correctly the first time, reducing bounce-around and wasted effort. It gives technicians the context and diagnostics they need up front, speeds up routine work without risking customer impact, and helps junior staff handle more tickets safely. The result is faster response, fewer mistakes, and clearer auditability for leadership.

Explanation:

Ingests new tickets from any PSA or helpdesk and enriches them with live endpoint, user, and tenant context

Classifies severity, category, and likely root cause using bounded AI reasoning

Automatically runs safe, non-disruptive diagnostics such as log collection, service checks, and backup validation

Stops before any disruptive remediation and routes those actions for approval

Escalates to a human technician when confidence is low or separation-of-duties policies apply

What's in It for the MSP:

Tickets get routed correctly the first time, reducing bounce-around and wasted effort

Technicians start with context and diagnostics already done, cutting time-to-resolution

Automation speeds up routine work without risking customer impact

Junior techs handle more tickets safely, easing pressure on senior staff

Customers see faster responses and fewer mistakes, while leadership gets clear auditability

MSP Acquisition Integration Coworker

Integrates: PSA, RMM/UEM, Security (EDR), Backup, Identity, Documentation

Explanation:

Detects the acquisition of a new MSP and initiates a controlled integration workflow

Inventories tools, customers, technicians, and configurations from the acquired MSP

Maps differences in SOPs, security posture, patching standards, and access models
Safely migrates customers, endpoints, and technicians into the acquiring MSP's standards
Runs side-by-side validation to confirm monitoring, backup, and security are active
Requires approval for destructive changes or policy overrides
Logs every action and decision for audit, compliance, and post-merger review

What's in It for the MSP:

Acquisitions integrate faster without breaking customer service
Risk is reduced by standardizing security, patching, and access before changes go live
Operational consistency improves instead of inheriting chaos
Senior engineers aren't pulled into weeks of manual cleanup
Leadership gains clear visibility into progress, gaps, and readiness
The MSP turns M&A into a repeatable playbook instead of a one-off fire drill [\[NR1\]](#)

Disaster Recovery Orchestration Coworker

Integrates: Cove DR/Standby Image (incl. Azure recovery option), PSA, Approval Workflow, Customer Comms

Explanation:


Detects an outage (ticket, monitoring event, or customer call) and initiates a DR workflow.

Enforces step-level approvals for failover, network changes, and "declare disaster" moments, then executes only what's authorized.

Captures the full DR timeline: what was approved, what ran, where it recovered, and when service was restored.

Other use cases considered

Network Incident Resolution: Deterministic + Non-Deterministic in Action



A network error fires. The **N-able Monitor Expert** catches it and triggers the customer's **AI Coworker**, which orchestrates the entire response (it has their SOPs and connections to their favored network vendor).

The Coworker calls the **Vendor Network Expert**, a third-party agent that pulls device state, topology, and error logs. That data flows to the **N-able Vulnerability Expert**, which analyzes it against known CVEs and determines a specific remediation: patch a firewall rule on Device X.

The Coworker then consults the **N-able Script Expert**, which searches the Script Repository and finds a validated script that performs exactly this action via direct API calls. The Script Expert flags it as deterministic, meaning no model reasoning is required to execute it.

The Coworker hits a decision gate: "Can this action be completed without a model?" The answer is yes. So instead of having the LLM generate API calls (burning tokens and risking hallucinated parameters), the Coworker hands off execution to the Script Execution Engine, which runs the script directly. The agent process pauses. The script runs, the result comes back, and the Coworker logs the outcome.

Five agents collaborated using model reasoning to diagnose the problem and identify the fix. But the moment the action could be performed deterministically, the system stopped spending tokens and handed off to validated, auditable code.

Monetization: The ACP as a Revenue Architecture

Monetization: The ACP as a Revenue Architecture

The Per-Device Model Is Broken

Per-device pricing worked when a scan surfaced 10 findings a day. AI agents surface 3,000 in one pass. The volume of valuable work exploded and the MSP currently does charge more for the same 500 laptops because AI found more things to fix.

This is industry-wide. SAP is shifting to consumption-based AI pricing. McKinsey found the TAM for AI+SaaS now includes labor itself, and incumbents must price on consumption to capture it. BCG found 40% of enterprise buyers cite seat reduction as their primary lever to cut software spending. [[si SAPinsider](#)] [[mckinsey.com](#)] [[scm Strategic Management Consulting](#) | [Boston Consulting Group](#)]

If the pricing model counts laptops while agents do 100x the work, revenue stays flat while value multiplies.

How the Market Already Prices This

n8n charges per workflow execution. €20/month for 2,500 runs. No per-user or per-workflow caps. [[🇺🇳 Latenode — AI Workflow Automation with 1,200+ apps and LLMs Built In](#)]

[🔗 AI Workflow Automation Software & Tools](#) charges per operation. Five-step workflow, five operations. Shifted to credits in 2025 to handle AI workloads. [[🔗 AI Workflow Automation Software & Tools](#)]

Salesforce Agentforce launched at \$2/conversation, then Flex Credits at \$0.10/action. Three buying models: pay-as-you-go, pre-commit, pre-purchase. [[🔗 Salesforce: The #1 Agentic AI CRM](#)]

Zendesk charges \$1.50 per automated resolution. If the AI escalates to a human, you pay nothing. [[🔗 AI-Powered Service Platform | Zendesk](#)]

Same conclusion everywhere: price the work, not the laptop.

Seven Revenue Layers (They Stack)

1. **Platform Access (Subscription)** Access to the control plane, builder, audit system, key service. You pay for the right to build and run.
2. **Workflow Execution (Consumption)** Every Coworker run is a billable event. 200 tenants, 200 executions. The n8n model applied to MSP operations.
3. **Expert Licensing (Catalogue Revenue)** Each Expert carries its own recurring license. The Identity Expert, the Vulnerability Expert, the Halo Ticket Expert. MSPs activate what they need.
4. **Vendor Ecosystem (Marketplace Economics)** Three streams:
 - **Revenue share on Vendor Expert transactions.** 0% on the first \$X0K (seeds the ecosystem), 15% to \$5M, 10% above (rewards scale). Same logic as Shopify. [[🇺🇳 Shopify: The All-in-One Commerce Platform for Businesses - Shopify](#)]
 - **Certification fees.** One-time \$500 to \$2,000 registration. Waived for Technology Alliance partners.
 - **Ecosystem-driven execution volume.** Every Vendor Expert in the catalogue gives MSPs a reason to build another Coworker, which generates more executions (Layer 2), more tool transactions (Layer 5), and more inference calls (Layer 6). The vendor ecosystem amplifies every other layer.
5. **MCP and Tool Transactions (Metered Access)** Every MCP call, API call, or script execution can be metered. [[🔗 AI Workflow Automation Software & Tools](#)]'s per-operation model at the

tool level. Pricing reflects actual resource consumption.

6. **LLM Inference Pass-Through (Cost Recovery + Margin)** Bedrock customers pay inference cost plus margin. BYOK/BYOM customers skip this layer. The Deterministic Substitution Engine reduces N-able's cost every time it routes to a script instead of an LLM call. The execution still counts. The Substitution Engine is a margin engine.
7. **Outcome-Based Pricing (Future State)** Per resolved incident. Per onboarded employee. Per QBR generated. Requires clean baselines and attribution. Highest margins because price anchors to business value. [[z AI-Powered Service Platform | Zendesk](#)]

How the Layers Compound

An MSP subscribes (Layer 1). Activates three N-able Experts and two Vendor Experts (Layers 3, 4). Builds a Coworker that runs across 150 tenants nightly (Layer 2). Each execution generates MCP transactions (Layer 5) and inference calls (Layer 6). The vendor revenue share flows to N-able (Layer 4).

How the MSP Charges Their Customer

The ACP gives MSPs a new pricing vocabulary. Each Coworker is a named, demonstrable unit of value they can show the customer.

- **Bundled.** Include Coworkers in the monthly fee. Treat AI as a margin play.
- **Tiered.** Basic gets monitoring. Premium gets automated remediation plus QBR reports.
- **Consumption.** "Your environment ran 3,200 remediations this month. Here is the invoice."
- **Outcome-based.** Per onboarded employee. Per resolved incident. Capture the spread between execution cost and outcome value.

The (not named Marketplace) Vendor onboarding, access, and monetization process

The ACP vendor ecosystem follows the same structural pattern as every major platform marketplace: submission, review, certification, listing, ongoing compliance.

Onboarding a Vendors' AI Expert

A vendor submits four things:

1. **The Expert** — the deployable agent with its runbook and default configuration
2. **Declared MCP scopes** — what it accesses, what methods it calls, maximum credential scope
3. **Proof that the AI Expert passes our minimum test suite** (we will supply our working test environment, Vendor's team can install it and run it.)

4. **A sales contract, support contact and security response SLA**

Certification

N-able runs automated scope validation (does the Expert only do what it says it does?).

Type	Timeline
New Expert	??
Minor update (no scope change)	??
Scope change or new MCP access	??
Critical security patch	48 hours

Upon certification we ‘publish’ the AI Expert as an available resource in the marketplace.

Experts may set a list price for utilization (see [☰ Control Plane / Headless IT Tech in the AI Agent world | \[hardBreak\]\[hardBreak\]Monetization: The ACP as a Revenue Architecture](#))

Revenue Model

Keep it simple.

Tier	N-able Take
First \$50K annual revenue	0%
\$50K to \$5M	15%
Above \$5M	10%

One-time registration fee (\$500 to \$2,000) to cover review costs. Waived for existing Technology Alliance partners. Vendor retains full IP. N-able holds a distribution license.

Versioning

Updates enter certification review before publishing. New versions publish alongside old versions. MSPs pin to the version they built against. They adopt new versions when they choose. Security patches are the one exception: N-able can mandate adoption within 72 hours.

Breaking changes require the vendor to declare them upfront. Affected MSPs get notified with a specific impact assessment. Minimum 90-day migration window before the old version sunsets.

The Marketplace

MSPs discover Experts by category, vendor, product line compatibility, and rating. Each listing shows a star rating, number of MSPs using it, execution success rate, and which Experts it is commonly paired with.

Vendors see aggregate adoption, usage, failure rates, and revenue. They do not see individual MSP names or customer data.

Ongoing Governance

Three mechanisms:

- **Runtime monitoring** — the audit layer tracks whether the Expert stays within its certified scope during production use
- **Annual re-certification** — lighter than initial review, vendor confirms behavior, N-able runs automated validation
- **Emergency delisting** — if a critical vulnerability or data violation is found, N-able pulls the listing immediately and notifies affected MSPs