

Sources – Reading Guide

Sources – Reading Guide (set up to read for success)

How to use this: the **scout** files on disk are your reading copies – each is one-claim-per-line with a [CS:] confidence tag and a clickable source URL. You don't have to trust my synthesis; every claim points back to where it came from. Read in the order below; trust the tags.

* Start here – the 4 that carry the whole argument

- **Nicole's demand study** (the reason this exists) – scout/confluence-blind-interviews-study.md → live page: <https://n-able.atlassian.net/wiki/spaces/AI/pages/62064099623> (needs your n-able login)
- **Competitive landscape** (who's doing it / in pieces) – scout/market-competitive-landscape.md
- **TAM + growth** (is the market real/growing) – scout/market-tam-and-growth.md
- **N-Zo behavioral signal** (time savings + trust gaps) – scout/condens-nzo-and-preview.md

Read-for-the-question map

Your question	Read this file	Live source
Is it wanted? Who wants it?	confluence-blind-interviews-study.md	Confluence study (login)
Are the time-savings / trust-gaps real?	condens-nzo-and-preview.md	3 Condens links (below)
Who's already in market? Who does pieces well?	market-competitive-landscape.md	–
How big / is it growing?	market-tam-and-growth.md	–
How do MSPs actually <i>run</i> a QBR?	web-qbr-meeting-ethnography.md	–
What's at stake emotionally / what fails?	web-qbr-human-factors.md	–
Are the brief's claims/competitors/vendors even real?	web-competitive-intel.md · web-integration-reality-check.md · web-benchmarks-and-sources.md	–

[LOCKED] Live links (need access – open in a browser)

- **Confluence – BLIND INTERVIEWS ACP study:**
<https://n-able.atlassian.net/wiki/spaces/AI/pages/62064099623> (needs n-able SSO)
- **Condens – N-Zo perceived time-saving:**
<https://share.condens.io/aw4c8940rw5nxGD1tfxrC> (JS-gated – WebFetch couldn't read it; open in a browser to get the exact 90→17.5min figures the brief cites)
- **Condens – N-Zo in preview (trust/export gaps):**
<https://share.condens.io/aotjLmXXEDnyyWKh8lfdF>
- **Condens – Empower 2026 (exec-narrative demand):**
<https://share.condens.io/a4eNwmzMalaR3Qqg9uByA>

Web/trade sources are open – every URL is inline in the scout files; just click.

Trust ladder – read with the right skepticism

- **Strongest (bank it):** Nicole's study (internal primary, but see caveat) · competitor **funding/pricing** (independent secondary, [CS: HIGH]).
- **Directional only:** the demand ranking ("#1, total 31") – it's **structured qual from n=8**, not projectable quant. Read it as *what 8 experts prioritized*, not *what the market will do*.
- **[WOBBLY]** – use for spread, not point facts: market-size reports (firms disagree ~50%: \$309B-\$460B) · all QBR trade-press (vendor blogs, no independent study).
- **!! DO NOT CITE (traps):**
 - "82% / 85% QBR retention stats" – don't trace to any primary source.
 - "~30% AI efficiency" (VC rollups) – FABRICATION RISK, no source.
 - Kenneth "\$300+/user/mo" & Vitali "15-20 hrs" – respondent-reported, unverifiable (public QBR-prep is 3-8 hrs).
 - ScalePad "displacement window" – public record contradicts (Lifecycle Manager X is the *new* product being sold).

How to read a scout file (30 seconds)

Each line = one claim + a [CS:] tag + a URL.

- VERIFIED / HIGH → trust it.
- MEDIUM → plausible, sanity-check.
- LOW / **[WOBBLY]** → directional; don't quote as fact.
- UNKNOWN / **[?]** → a named gap, not a finding.
- FABRICATION RISK → skip.

The tag tells you how hard to lean. That's the whole discipline – fluency ≠ accuracy.

All files live in prototypes/acp-qbr-analyst/scout/ (in the handoff packet) or scout_input/acp-qbr-analyst/ (working copy). Same content.